

# Web Search Basics and Web Crawling

## CS 7263 Information Retrieval Lecture 01

Jiho Noh

Department of Computer Science  
Kennesaw State University

Fall 2025



- 1 **The Web and Search Engines**
  - History of Search Engines
  - From Keyword to GenAI Search
- 2 **Fundamentals of Web Crawlers**
  - Web as a Graph
  - Coverage, Freshness, Politeness
  - Handling Duplicates

# The World Wide Web

- The Web is a system of interlinked hypertext documents accessed via the Internet.
- Developed by Tim Berners-Lee in 1989 at CERN
  - ▶ To organize research documents for scientists.
  - ▶ Combined the idea of *hypertext* to link documents with the documents available via FTP.
  - ▶ Developed HTTP network protocol, URLs, HTML, and the first web browser.

# Early Search Engines

**1990** Alan E. Hall's **Archie** was the first search engine

- ▶ Indexes FTP sites, not the web
- ▶ Searchable database of file names, allowing regular expressions (regex) search

**1993** **Gopher** was the first search engine to index the web; searching names of text files

**1993** Spider robots were developed to crawl the web and collect URL's

**1994** Stanford grad students David Filo and Jerry Yang created **Yahoo!**

- ▶ A directory (hierarchy) of web pages, not a search engine
- ▶ Manually collecting popular web pages
- ▶ Users could submit their own pages to be indexed

# Directory of Web Pages

## [Yahoo](#) - A Guide to WWW

[ [What's New?](#) | [What's Cool?](#) | [What's Popular?](#) | [Stats](#) | [A Random Link](#) ]

[Y](#) [Top](#) [Up](#) [Search](#) [Mail](#) [Add](#) [Help](#)

- [Art](#)(466) NEW
- [Business](#)(6426) NEW
- [Computers](#)(2609) NEW
- [Economy](#)(743) NEW
- [Education](#)(1487) NEW
- [Entertainment](#)(6199) NEW
- [Environment and Nature](#)(193) NEW
- [Events](#)(53) NEW
- [Government](#)(1031) NEW
- [Health](#)(367) NEW
- [Humanities](#)(163) NEW
- [Law](#)(163) NEW
- [News](#)(185)
- [Politics](#)(148) NEW
- [Reference](#)(474) NEW
- [Regional Information](#)(2606) NEW
- [Science](#)(2634) NEW
- [Social Science](#)(93) NEW
- [Society and Culture](#)(648) NEW

23836 entries in *Yahoo* [ [Yahoo](#) | [Up](#) | [Search](#) | [Mail](#) | [Add](#) | [Help](#) ]

[yahoo@akebono.stanford.edu](mailto:yahoo@akebono.stanford.edu)

Copyright © 1994 David Filo and Jerry Yang

# Early Search Engines (Cont.)

**1994** Brian Pinkerton developed WebCrawler at the University of Washington

- ▶ Eventually became part of Excite and AOL

**1998** Larry Page and Sergey Brin developed **BackRub** (Google's original name) at Stanford

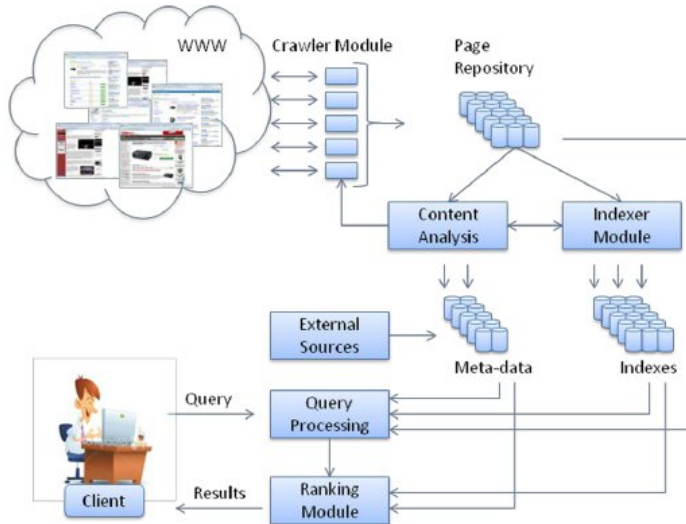
- ▶ A search engine that used *PageRank* algorithm to rank web pages

**1998** Microsoft launched **MSN Search** (now Bing)

- ▶ A web search engine that used a combination of algorithms and human editors to rank web pages
- ▶ Bing now evolved focused on features like image and video search, and providing contextual search suggestions

# Core Components of Traditional Search Engines

- **Crawling**
  - ▶ Discovering and fetching web pages from the internet
- **Indexing**
  - ▶ Organizing and storing the crawled data in a way that allows for efficient retrieval
- **Query Processing**
  - ▶ Analyzing and interpreting user queries to retrieve relevant results
- **Ranking**
  - ▶ Determining the order in which search results are presented to the user



**Figure:** The Architecture of a Web Search Engine; image by Michalis Vazigiannis



# Challenges of Web Search

- **Scale of the Web**

- ▶ Large volume of data
- ▶ Distributed nature: Documents spread across millions of web servers

- **Language understanding**

- ▶ Language Ambiguity: Words can have multiple meanings, and the same meaning can be expressed in different ways
- ▶ Unstructured data: Most web content is unstructured, making it difficult to extract relevant information

- **Heterogeneous content**

- ▶ Different formats and media types: Web pages can be in HTML, PDF, images, videos, etc.
- ▶ Different languages: The web is multilingual, and search engines need to handle multiple languages

- 1 **The Web and Search Engines**
  - History of Search Engines
  - From Keyword to GenAI Search

- 2 **Fundamentals of Web Crawlers**
  - Web as a Graph
  - Coverage, Freshness, Politeness
  - Handling Duplicates

# Limitations of keyword search

- **Keyword Search**

- ▶ Users enter keywords, and the search engine retrieves documents containing those keywords
- ▶ Limited understanding of user intent and context
- ▶ Often returns irrelevant or low-quality results

# Rise of Conversational AI and LLMs

- **Conversational AI**

- ▶ AI systems that can engage in natural language conversations with users
- ▶ Examples: Chatbots, virtual assistants (e.g., Siri, Alexa)

- **Large Language Models (LLMs)**

- ▶ Deep learning models trained on vast amounts of text data
- ▶ Capable of understanding and generating human-like text
- ▶ Examples: GPT-3, BERT, T5

# Case Study: Perplexity AI

The screenshot displays the Perplexity AI search interface. At the top, the search query is "What are the best places to visit in Tokyo". The interface includes a sidebar with navigation options: Home, Discover, and Spaces. The main content area shows the search results under the "Perplexity" tab. It features three source cards: TripAdvisor's "THE 15 BEST Things to Do in Tokyo (2025) - Must-See...", gotokyo's "Top tourist attractions in Tokyo: the best sightseeing spots", and cntraveler's "28 Best Things to Do in Tokyo - Condé Nast Traveler". Below these cards is a carousel of four images representing various Tokyo locations. The text "Here are some of the best places to visit in Tokyo:" is followed by a list item for Shinjuku Gyoen National Garden, describing its serene lakes, diverse landscapes, and cherry blossoms.

What are the best places to visit in Tokyo

Perplexity Images Sources 8 1 task ↗

Tripadvisor THE 15 BEST Things to Do in Tokyo (2025) - Must-See...

gotokyo Top tourist attractions in Tokyo: the best sightseeing spots

cntraveler 28 Best Things to Do in Tokyo - Condé Nast Traveler +5 sources

Here are some of the best places to visit in Tokyo:

- **Shinjuku Gyoen National Garden** This tranquil garden in Shinjuku features serene lakes and diverse landscapes, including Japanese, English, and French styles, making it ideal for scenic strolls and picnics [1]. The park is especially beautiful when the cherry blossoms are blooming

# How Perplexity AI Works

Feature	Traditional Search Engine	Perplexity (GenAI Search)
Query Input	Keyword-based	Conversational, natural language
Results	Ranked list of links	Direct answers with citations
Indexing	Inverted index, static pages	Real-time web search, dynamic
Personalization	Limited, based on history	Context-aware, remembers threads
Advanced Capabilities	Boolean, phrase search	Summarization, file analysis
Transparency	Limited	Source citations in every answer

**Figure:** Perplexity generated answer to “GenAI-based search engines”

- 1 **The Web and Search Engines**
  - History of Search Engines
  - From Keyword to GenAI Search
- 2 **Fundamentals of Web Crawlers**
  - Web as a Graph
  - Coverage, Freshness, Politeness
  - Handling Duplicates

# Web Crawling

- **Web crawling** is the process of locating, fetching, and storing the pages available in the Web
- A program that automatically downloads web pages is called a *web crawler*
- Follow all links on these visited pages recursively to find additional pages.
- Other names of web crawlers are:
  - ▶ spiders
  - ▶ scrapers
  - ▶ robots
  - ▶ harvesters



# Retrieving Web Pages

- Each web page has its own unique *Uniform Resource Locator* (URL)
- Most web pages can be retrieved via *Hypertext Transfer Protocol* (HTTP)


**https**://**ccse.kennesaw.edu**/**cs/course-descriptions.php**

**scheme**                      **hostname**                      **resource**

**Figure:** A uniform resource locator (URL)




# Simple HTTP Transaction (cont.)



- **(2). Connect:** Client establishes TCP connection with the IP address
- Client sends SYN packet
- Web server sends SYN-ACK packet
- Client sends ACK packet, concluding the *three-way handshake* for TCP connection

# Simple HTTP Transaction (cont.)



- (3). **Send**: Client sends the HTTP request.
- (4). **Wait**: Client waits for the server to respond to the request.
- (5). **Load**: Client keeps communicating with the server to load all the TCP segments.
- Client sends a a FIN packet to close the TCP connection.

# HTTP Requests

An *HTTP Request* is made by a client to a web server for accessing a resource on the server.

## HTTP request components

- A request line (<method><url><HTTP version>)
- A series of HTTP headers
- A message body, if needed

# What's going on “under the hood” using cURL

- cURL is a command-line tool for transferring data with URLs
- With `-verbose` option, cURL shows the details of the HTTP request and response
- The printout contains four components:
  - ▶ Additional info. such as http transaction info, provided by curl (prefixed `*`)
  - ▶ “header data” sent by curl (prefixed `>`)
  - ▶ “header data” received from the server (prefixed `<`)

```
» curl -v -s -stderr https://www.google.com
```

```

* [HTTP/2] [1] [:authority: www.google.com]
* [HTTP/2] [1] [:path: /]
* [HTTP/2] [1] [user-agent: curl/8.7.1]
* [HTTP/2] [1] [accept: */*]
> GET / HTTP/2 Request Line
> Host: www.google.com
> User-Agent: curl/8.7.1 Request Headers
> Accept: */*
>
* Request completely sent off
< HTTP/2 200
< date: Sun, 27 Apr 2025 21:26:20 GMT
< expires: -1 Response Code
< cache-control: private, max-age=0
< content-type: text/html; charset=ISO-8859-1
< content-security-policy-report-only: object-src 'none';base-uri 'self';script-src 'nonce-Pi
https: http;;report-uri https://csp.withgoogle.com/csp/gws/other-hp
< accept-ch: Sec-CH-Prefers-Color-Scheme
< p3p: CP="This is not a P3P policy! See g.co/p3phelp for more info."
< server: gws
< x-xss-protection: 0
< x-frame-options: SAMEORIGIN
< set-cookie: AEC=AVcja2eR6_h5T8hN7bjQR8K61Muro-cb9bH56ggt8MaH3zCuzAjXtOmrA; expires=Fri, 2
< set-cookie: NID=523=W56aHRRqze6DmSSB3wEBR-3ZRLh1lvoYYVHg1fdzBJCUh_vvEK1QiQmEpFZ9okowyhEGIdr
D2GM_Som-YqP6D1_Ev1d2YVegPfp4_rYz6-JdRl2iEhCLx_ImYfv89PeD4U8; expires=Mon, 27-Oct-2025 21:26
< alt-svc: h3=":443"; ma=2592000,h3-29=":443"; ma=2592000
< accept-ranges: none
< vary: Accept-Encoding
Response Body
<!doctype html><html itemscope="" itemtype="http://schema.org/WebPage" lang="en"><head><meta
. Google has many special features to help you find exactly what you're looking for." name="
UTF-8" http-equiv="Content-Type"><meta content="/images/branding/googleg/1x/googleg_standard
EWv0pncA">(function(){var _g={kET:'_KA0aTa4E-iNwhkP-qeWoAU',kEXPT:'0',202791,27,3497447,1119,

```

# HTTP Request Methods


Method	Description
GET	Requests a representation of the specified resource. Requests using GET should only retrieve data.
HEAD	Asks for a response identical to a GET request, but without the response body.
POST	Submits an entity to the specified resource, often causing a change in state or side effects on the server. (e.g., submitting form inputs)
PUT	Replaces all current representations of the target resource with the request payload. (e.g., saving an object)
DELETE	Deletes the specified resource

- [Full list of HTTP Methods](#)
- [Full list of HTTP Response Stat codes](#)




# Web Graph


Web crawlers exploit the  
hyperlink structure of the Web



# DFS vs. BFS



Depth-first search



Breadth-first search

# DFS vs. BFS

- **Depth-First Search (DFS)**

- ▶ Requires less memory
- ▶ A seed page may not get crawled in the end

- **Breadth-First Search (BFS)**

- ▶ Explores the Web uniformly outward
- ▶ Requires memory of all nodes on the previous level (exponential in depth)
- ▶ Can avoid too many requests to a single host

- **Both,**

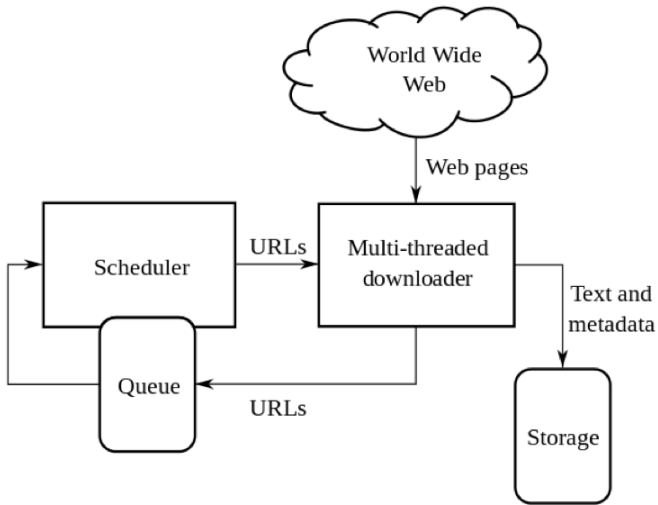
- ▶ Implemented using a queue of links (FIFO, LIFO)
- ▶ Heuristically ordering the Queue give a “focused crawler” that directs its search towards “interesting” pages.

# Web Crawling Process

A crawler maintains a **frontier** with a collection of pages to be crawled and iteratively fetches and stores pages from it.

- ➊ Initialize the frontier with a list of seed pages
- ➋ Selects the next page carefully, with the politeness and performance in considerations
- ➌ Downloads the selected pages, and extracts the hyperlinks within
- ➍ Locates new pages by parsing the downloaded pages
- ➎ Processes the new URLs and filters before being added to the frontier
- ➏ Continues downloading until the fetch queue gets empty or a satisfactory number of pages are downloaded.

# Web Crawling Process

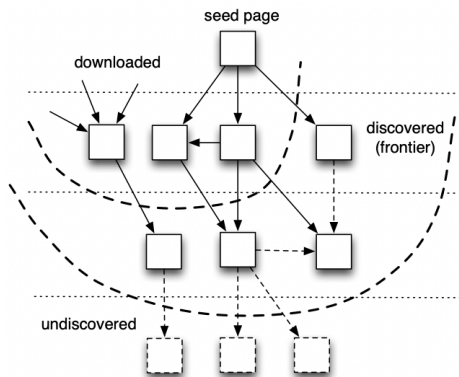


# Web Crawling — Frontier

Crawler frontier is a data structure used to store eligible URLs for crawling. It can be seen as a priority queue.

Crawler divides the Web into three sets

- downloaded
- discovered
- undiscovered



# URL Prioritization

A state-of-the-art web crawler maintains two separate queues for prioritizing the download of URLs:

## Discovery queue

- Downloads pages pointed by already discovered links
- Tries to increase the **coverage** of the crawler

## Refreshing queue

- re-downloads already downloaded pages
- tries to increase the **freshness** of the repository

- 1 **The Web and Search Engines**
  - History of Search Engines
  - From Keyword to GenAI Search
- 2 **Fundamentals of Web Crawlers**
  - Web as a Graph
  - Coverage, Freshness, Politeness
  - Handling Duplicates



# Goals of Web Crawlers

- **Coverage**

- ▶ No crawler can index all the Internet content,
- ▶ Hence it should prioritize high-quality content
- ▶ Discovering new pages and web sites as they appear online

- **Freshness**

- ▶ The recency of the content in your index
- ▶ A crawler should measure the rate at which each page changes
- ▶ A crawler should balance between the coverage and freshness goals

- **Politeness**

- ▶ Be respectful of other's web service resources by obeying the requested policies
- ▶ Put a delay between a consecutive downloads from the same server

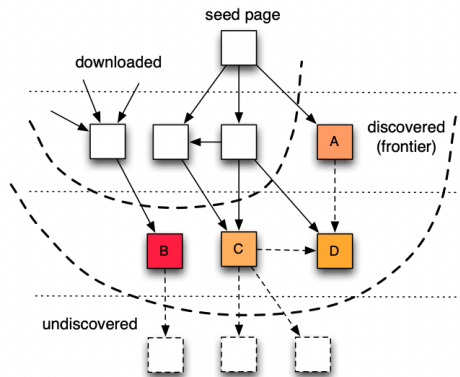
# URL Selection Strategies

Methods to select the next page to crawl

- **Breadth-First Search (BFS):** This distributes requests across domains relatively well and tends to download high-PageRank pages early
- **Backlink count:** Prioritize pages with more in-links from already-crawled pages.
- **Larger sites first:** Prioritize pages on domains with many pages in the frontier.
- **Partial PageRank:** Approximate PageRank scores are calculated based on already-crawled pages.

\* **PageRank** measures the importance of a webpage by link analysis

# URL Selection Strategies (Discovery)




*Which webpage will be selected for next download?*

- Random: Any of (A, B, C, D)
- Breadth-first: A
- In-degree: C
- PageRank: B

\* More intense red color indicates higher PageRank score

# DNS Caching


- Before a web page is crawled, the host name needs to be resolved to an IP address
- Since the same host name appears many times, DNS entries are locally cached by the crawler



# Freshness vs. Age

- Webpage modification rate has a huge variations.
- Crawlers need to measure the rate at which each page changes.
- Two metrics:
  - ▶ **Freshness:** a page is *fresh* if the crawl has the most recent copy of a web page
  - ▶ **Age:** a page has age 0 until it is changed and then its age grows

# Freshness vs Age



**Freshness is binary, age is continuous.**

- It's better to re-crawl pages when the age of the last crawled version exceeds some limit.
- The age of a page is the elapsed time since the first update after the most

# Expected Page Age

Suppose we have a page with change frequency  $\lambda$ , (i.e., a page changes  $\lambda$  times in a one-day period)


The expected age of a page  $t$  days after it was crawled depends on its update probability:

$$age(\lambda, t) = \int_0^t P(\text{page changed at time } x)(t - x)dx$$

On average, page updates follow a Poisson distribution – the time until the next update is governed by an exponential distribution. This makes the expected age:

$$age(\lambda, t) = \int_0^t \lambda e^{-\lambda x}(t - x)dx$$

# Freshness vs Age



**Fig. 3.7.** Expected age of a page with mean change frequency  $\lambda = 1/7$  (one week)



# Politeness

- In general, we don't want to make heavy demands on other's web service resources.
- Some web content is considered private or under copyright.
- Politeness policies can mediate the conflicts between search providers and the web sites.

# Our Web



**Figure:** Web by indexability and accessibility.

# Politeness Policy — Obey Robot Exclusion

- A standard from the early days of the Web
- The owner of a website can specify that robots should not crawl/index certain areas.
- **visible  $\neq$  accessible  $\neq$  storable**
- Always check before crawling the website

# Robot Exclusion

## Robot Exclusion Protocol

- Two components:
  - ▶ **Robots Exclusion Protocol:** Site wide specification of excluded directories. Specified in a file called 'robots.txt'
  - ▶ **Robots META Tag:** Individual document tag to exclude indexing or following links

# Robots.txt

'Robots.txt' is a file contains a list of excluded directories for given robot (user-agent)

## Exclude all robots from the entire site:

```
User-agent: *  
Disallow: /
```

## Exclude specific directories:

```
User-agent: *  
Disallow: /tmp/  
Disallow: /cgi-bin/  
Disallow: /users/paranoid/
```

## Exclude a specific robot:

```
User-agent: GoogleBot  
Disallow: /
```

# Politeness Policy — Robots.txt

## Robots.txt

```
user-agent: googlebot      # all services
disallow: /private/       # disallow this directory

user-agent: googlebot-news # only the news service
disallow: /               # on everything

user-agent: *              # all robots
disallow: /something/     # on this directory

user-agent: *              # all robots
crawl-delay: 10           # wait at least 10 seconds

disallow: /directory1/    # disallow this directory
allow: /directory1/myfile.html # allow a subdirectory

host: www.example.com     # use this mirror
```

google robots.txt

# Robots META Tag

- Include META tag in HEAD section of a specific HTML document.  
`<meta name="robots" content="noindex, nofollow">`
- Content value is a pair of values for two aspects:
  - ▶ **index** | **noindex**: Allow/disallow indexing of this page
  - ▶ **follow** | **nofollow**: Allow/disallow following links on this page

# Robots META Tag

- Special values:
  - ▶ **all** = index, follow
  - ▶ **none** = noindex, nofollow
- Examples:

```
<meta name="robots" content="noindex, follow">
```

```
<meta name="robots" content="index, nofollow">
```

```
<meta name="robots" content="none">
```



# Sitemaps

- In addition to robots.txt, which asks crawlers not to index certain content, a site can request that certain content be indexed by hosting a file at `/sitemap.xml`.
- Sitemaps can direct your crawler toward the most important content on the site, indicate when it has changed, etc.
- ex. Google's sitemaps

- 1 **The Web and Search Engines**
  - History of Search Engines
  - From Keyword to GenAI Search
- 2 **Fundamentals of Web Crawlers**
  - Web as a Graph
  - Coverage, Freshness, Politeness
  - Handling Duplicates

# Consolidating Duplicate URLs

## URL Canonicalization Process

- Many possible URLs can refer to the same resource.
- Many URL normalization rules are available:
  - ▶ some can make mistakes.
  - ▶ the order of rules can also affect the results

`http://example.com/some/../folder?id=1#anchor`  $\Rightarrow$   
`http://example.com/some/../folder`  
 $\Rightarrow$  `http://example.com/folder`

# URL Syntax

- A URL has the following syntax:  
`<scheme>://<authority><path>?<query>#<fragment>`
- An authority has the syntax  
`<host>:<port-number>`
- A query passes variable values from an HTML form and has the syntax:  
`<variable>=<value>&<variable>=<value>...`
- A fragment is also called a reference or a ref and is a pointer within the document to a point specified by an anchor tag of the form  
`<ANAME="<fragment">`

# URL Canonicalization Example

W3lib example (used in Scrapy package)

Example process: w3lib (used by the Python *Scrapy* package)

- 1 Sort query arguments, first by key, then by value  
<https://example.com/item?id=1231&cat=service&cat=product>
- 2 Percent encode paths; non-ASCII characters are percent-encoded using UTF-8 (RFC-3986)  
<http://www.example.com/r\u00e9sum\u00e9>
- 3 Percent encode query arguments; non-ASCII characters are percent-encoded using passed encoding (UTF-8 by default)
- 4 Normalize all spaces (in query arguments) '+' (plus symbol)  
<https://example.com/item?id=1231&cat=service&cat=europeancafe>

# URL Canonicalization Example

## Example process: w3lib (used by the Python *Scrapy* package)

- 5 Normalize percent encodings case (%2f → %2F)  
<http://www.example.com/r%c3%a9sum%c3%a9>
- 6 Remove query arguments with blank values  
<https://example.com/item?id=1231&cat=service&cat=europeancafe&loc=>
- 7 Remove fragments  
[https://example.com/book.php#second\\_title](https://example.com/book.php#second_title)

# Keeping Similar or Duplicate Pages

- Multiple device types
  - ▶ <https://example.com/news/koala-rampage>
  - ▶ <https://m.example.com/news/koala-rampage>
- Dynamic URLs (e.g., search parameters, session IDs)
  - ▶ <https://www.example.com/products?category=dresses&color=green>
- Different naming schemes for blog posts
  - ▶ <https://blog.example.com/dresses/green-dresses-are-awesome/>
  - ▶ <https://github.com/alexeygumirov/pandoc-beamer-how-to>
- Different network protocols
  - ▶ <http://example.com/green-dresses>
  - ▶ <https://example.com/green-dresses>
  - ▶ <http://www.example.com/green-dresses>

# Duplicate Detection

## Why?

- More than 30% of pages are (near-) duplicate documents on the Web
- Copies, mirror sites, versioning, plagiarism, spam, etc.
- Detecting duplicate documents is important for conserving limited crawling and indexing resources
- Also, an essential tool in natural language processing

## How?

- For exact duplicates, use **checksumming** techniques
- For near-duplicates, use a threshold value for some similarity measure between pairs of documents (e.g., **fingerprints**)




# Duplicate Detection

## Checksums vs. Hash functions

- *Hash (Checksum) functions* are mathematical functions that maps variable length data into a fixed value.
- Same concept with different purposes
  - ▶ **Checksum** is to detect errors (or accidental changes) in data, often in a simple and fast manner.
  - ▶ **Hash** is to speed up comparison of data using fingerprints or create a hash table.
- ASCII parity bit is a 1-bit hash function (i.e., even/odd parity)

# Hash functions (proof of concept)



# Fingerprints

## Process of generating fingerprints using n-grams

- ➊ A document is converted into a sequence of words; all punctuation and formatting is removed.
- ➋ Words are grouped into (possibly overlapping) n-grams, for some  $n$ .
- ➌ A subset of the n-grams are selected to represent the document.
- ➍ The selected n-grams are hashed, and the resulting hashes stored in an inverted index.
- ➎ Documents are compared based on the number of overlapping n-grams.

# Fingerprints (example)

## (a) Original text

Tropical fish include fish found in tropical environments around the world, including both freshwater and salt water species.

## (b) 3-grams

tropical fish include, fish include fish, include fish found, fish found in, found in tropical, in tropical environments, tropical environments around, environments around the, around the world, the world including, world including both, including both freshwater, both freshwater and, freshwater and salt, and salt water, salt water species

## (c) Hash values

938 664 463 822 492 798 78 969 143 236 913 908 694 553 870 779

## (d) Selected hash values using *Omod4*

664 492 236 908

# Summary

- Questions?
- Discussion?