

Attention Mechanism and Transformer

CS 7263 Information Retrieval Lecture 12

Jiho Noh

Department of Computer Science
Kennesaw State University

Fall 2025



Topics

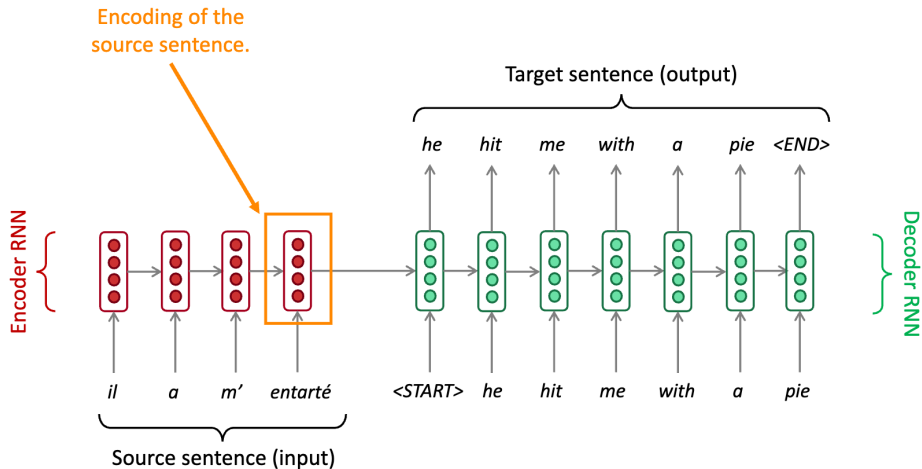
- 1 **Attention Mechanism**
- 2 Transformer Network
- 3 Pre-trained Language Models
 - BERT
 - GPT to ChatGPT

What is Attention Mechanism

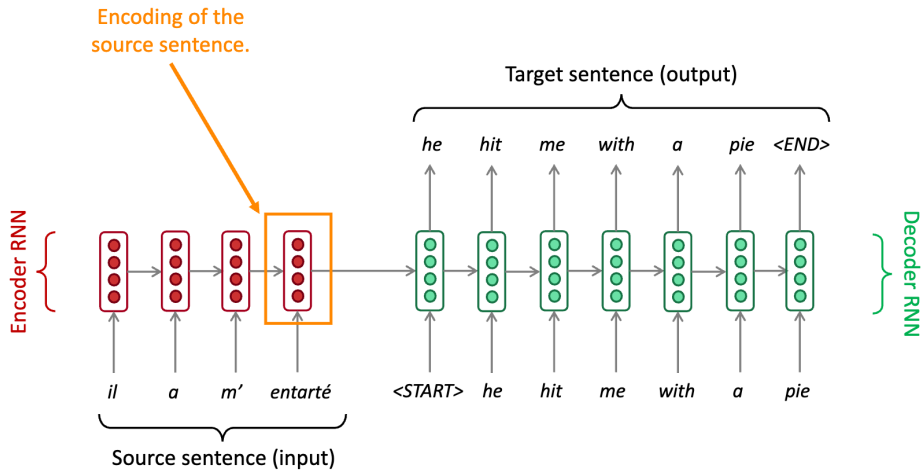
- Techniques used to **focus on specific parts of input data**.
- Attention **assigns weights** (or probabilities) to different parts of the input and selectively attend to the most relevant information.
- Commonly used in deep learning models, such as Transformer architecture, for natural language processing, image captioning, and other sequence-to-sequence models.
- Attention uses the notions of **Query, Key, Value**.

Attention Mechanism (Query, Key, Value)

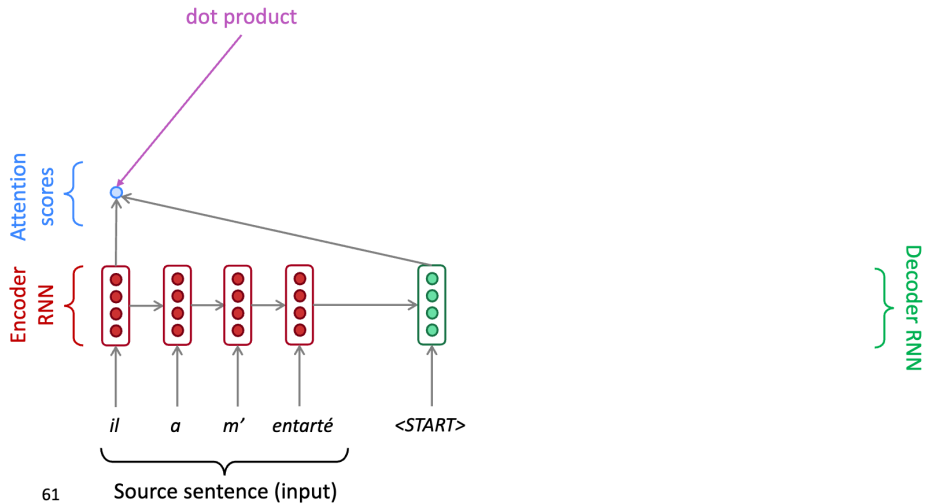
- What are the query, key, value vectors?
- These concepts come from retrieval systems. To make an analogy,
 - ▶ You type a **query** to search for some video on YouTube
 - ▶ The search engine will map your **query** against a set of **keys** associated with candidate videos in the database
 - ▶ Then, present you the best matched videos (**values**).

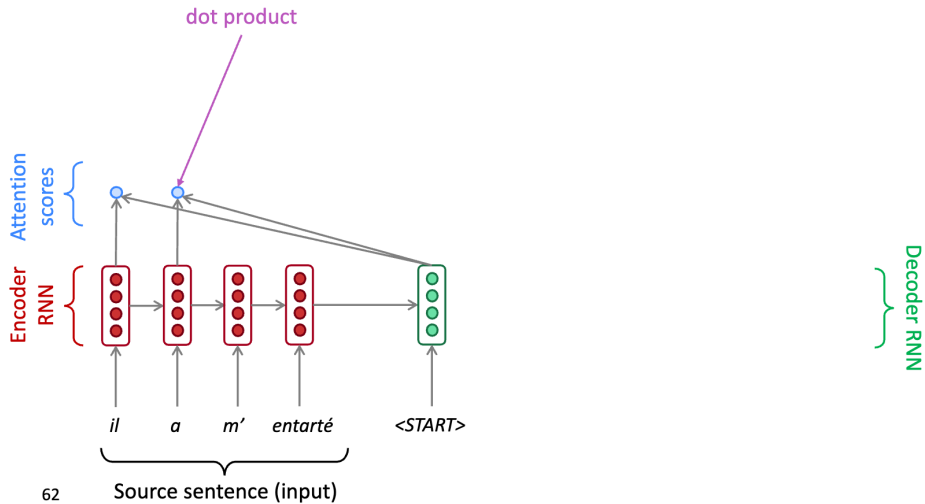


Problems with this architecture?

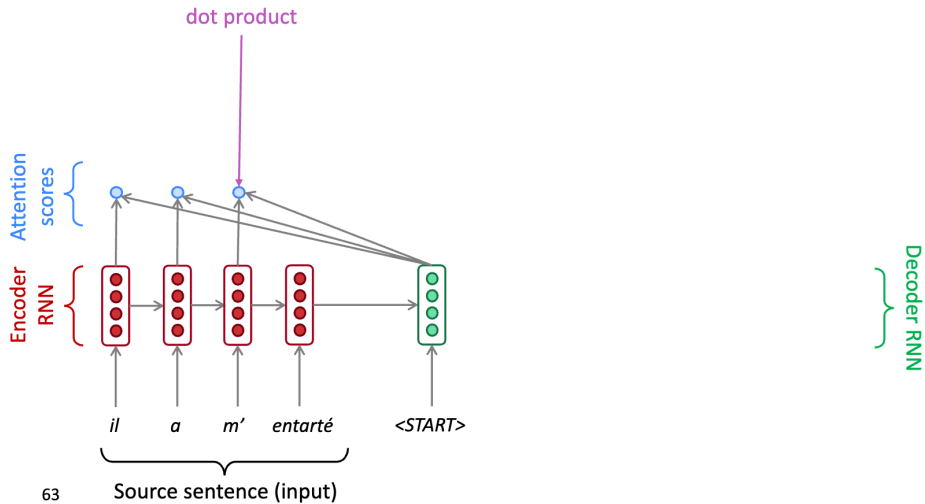


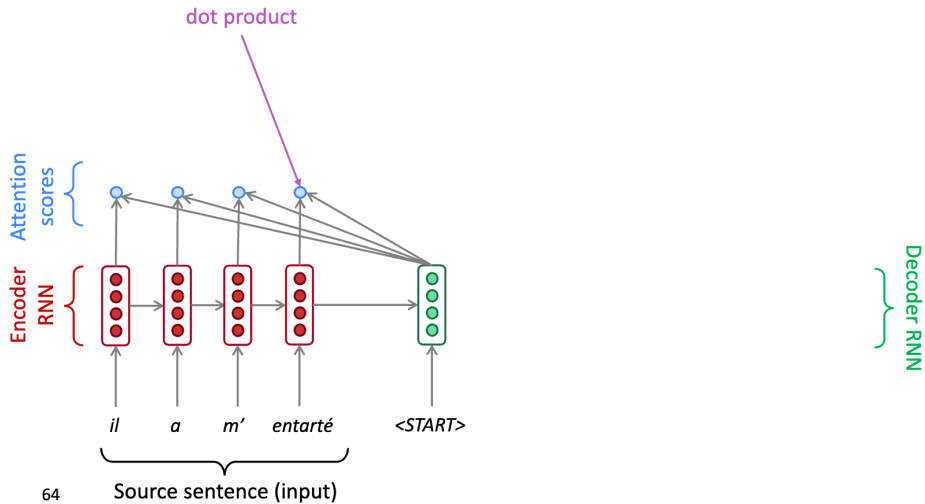
Problems with this architecture?



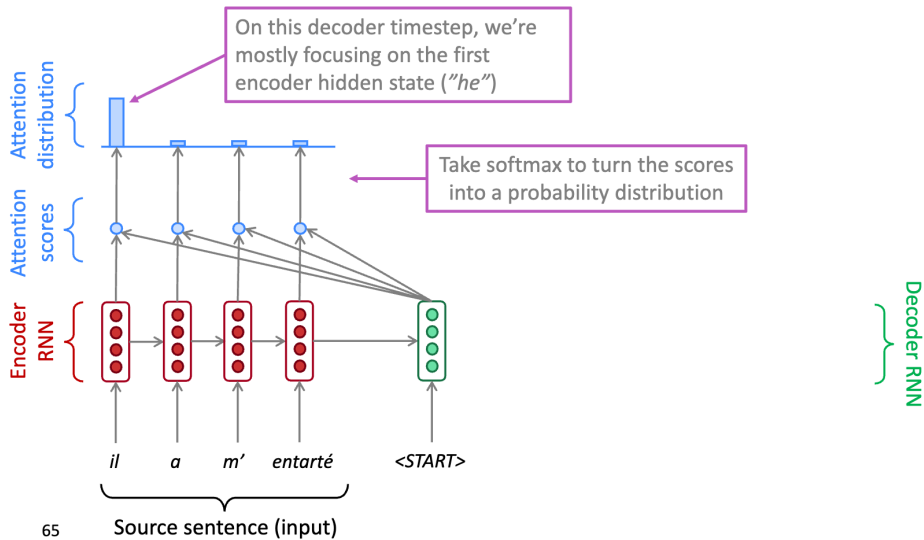


62

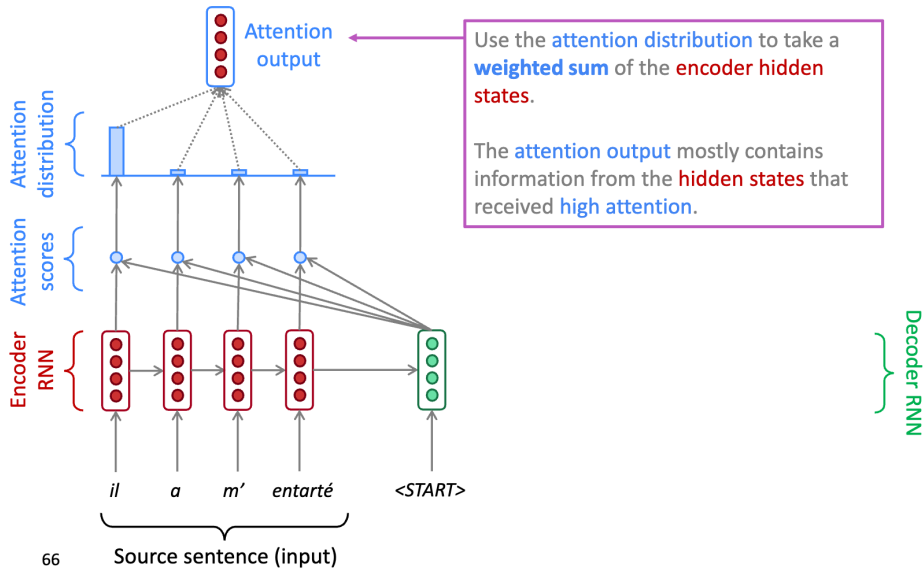


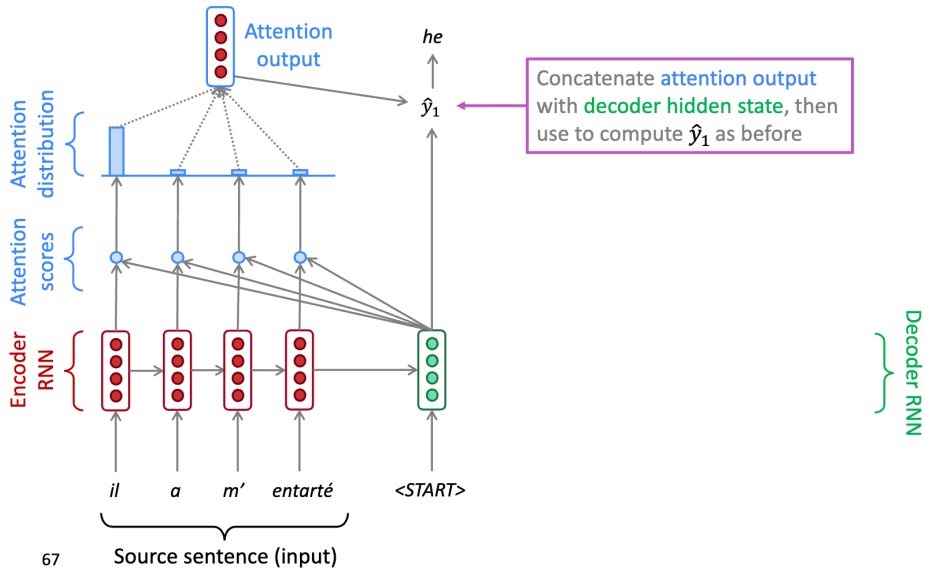


64

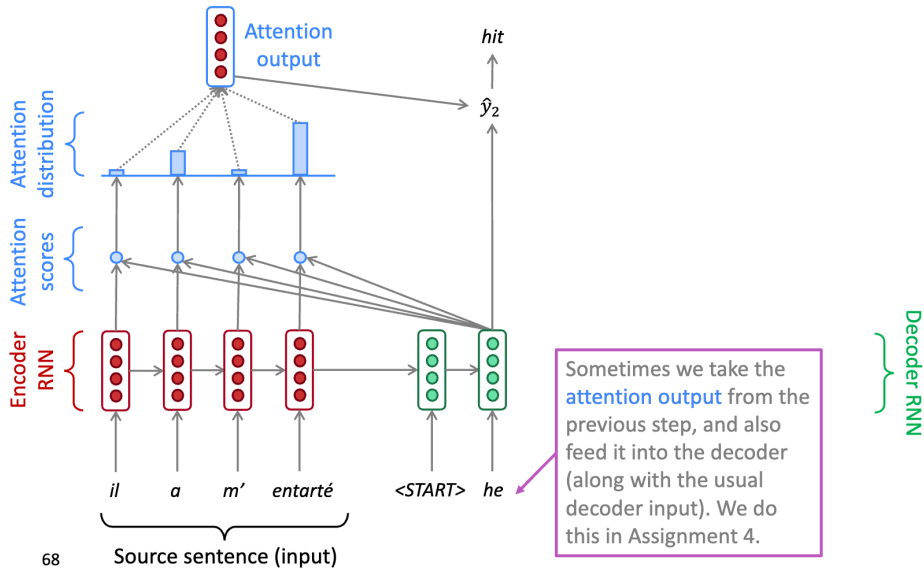


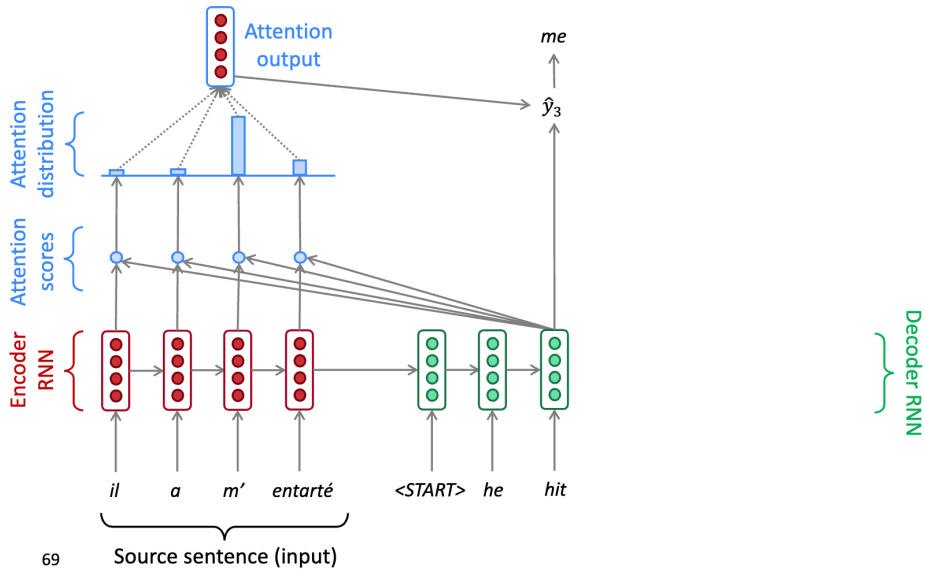
65

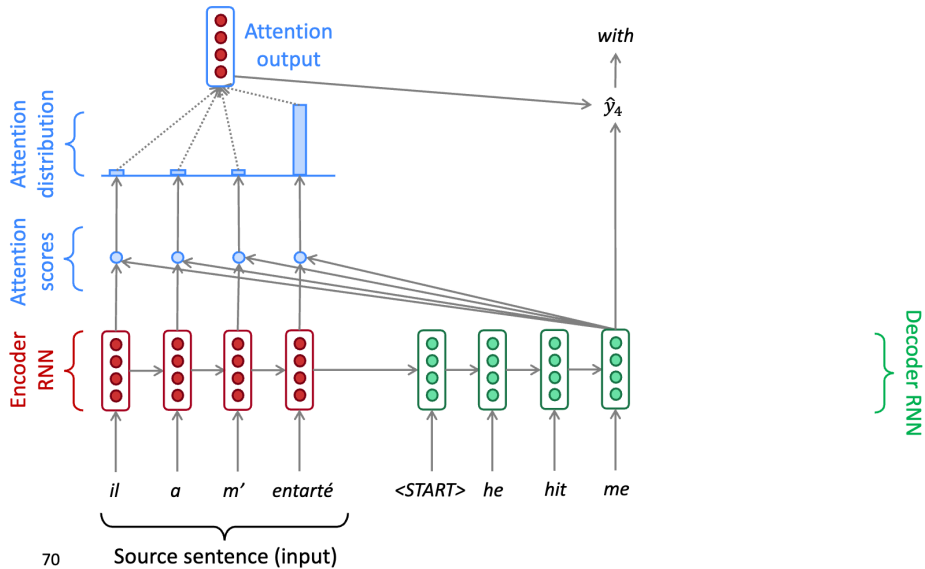




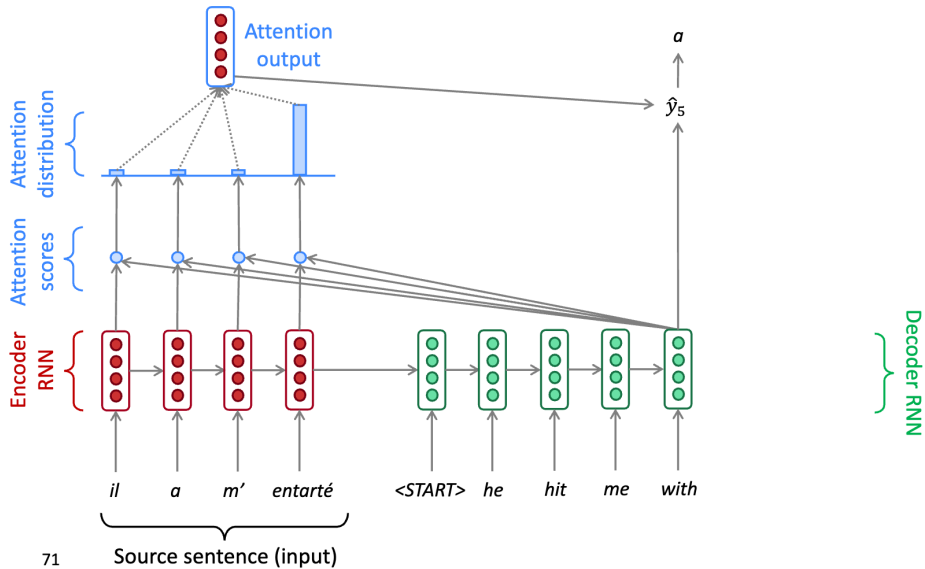
67

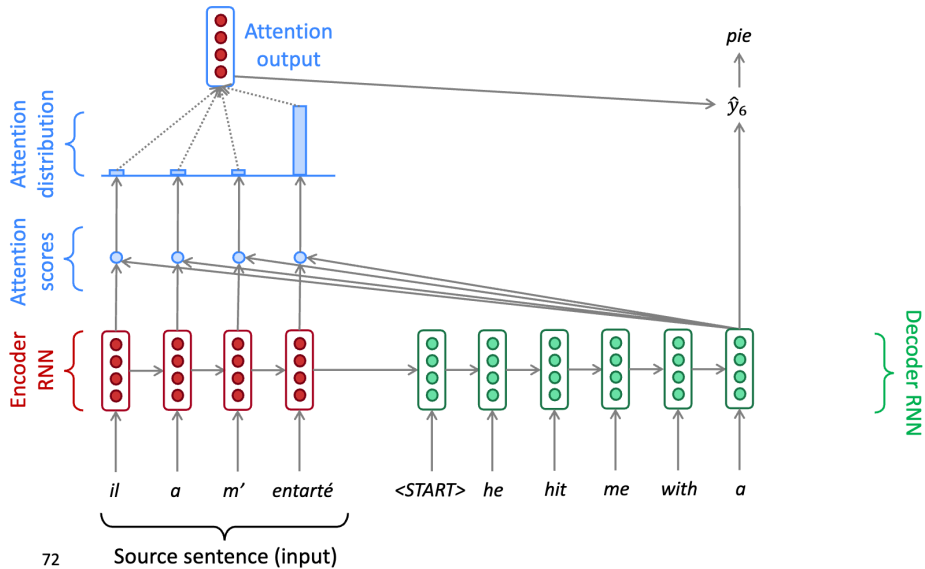






70





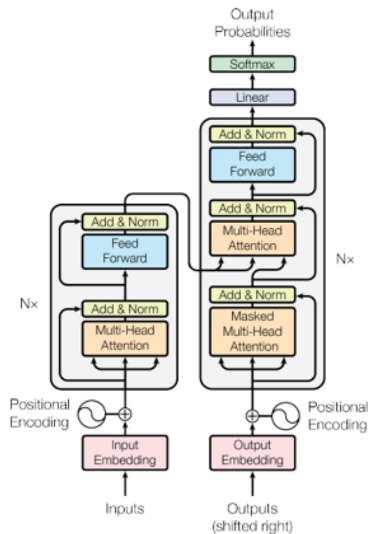
72

Advantages of Computing Attention

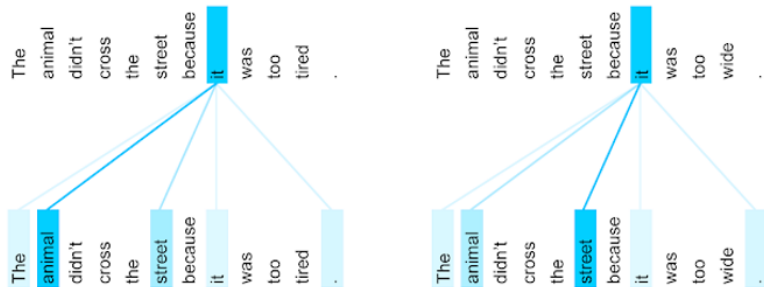
- Attention significantly **improves performance** (in many applications)
 - ▶ It's very useful to allow decoder to focus on certain parts of the source
- Attention **solves the bottleneck problem**
 - ▶ Attention allows decoder to look directly at source; bypass bottleneck
- Attention **helps with vanishing gradient problem**
 - ▶ Provides shortcut to faraway states
- Attention provides **some interpretability**
 - ▶ By inspecting attention distribution, we can see what the decoder was focusing on

Transformer

"Attention is All You Need"
(Vaswani et. al 2017)

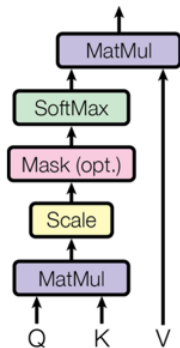


Self-Attention Example from NMT



- The encoder self-attention distribution for the word 'it' from the 5th to the 6th layer of a Transformer trained on English to French translation (one of eight attention heads)

Scaled Dot-Product Attention



Multi-Head Attention

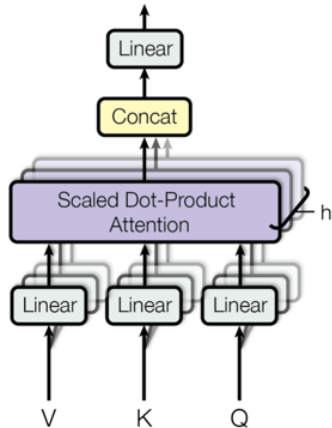


Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention consists of several attention layers running in parallel.

Topics

1 Attention Mechanism

2 **Transformer Network**

3 Pre-trained Language Models

- BERT
- GPT to ChatGPT

Transformer with Self-Attention Illustrated

“The Illustrated Transformer” by Jay Alammar

Self-Attention

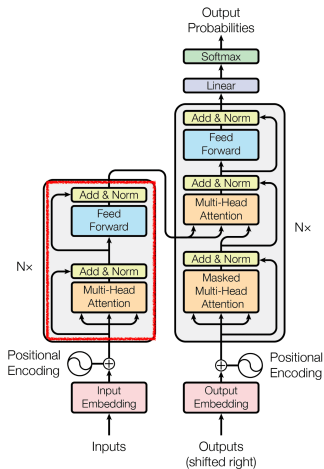
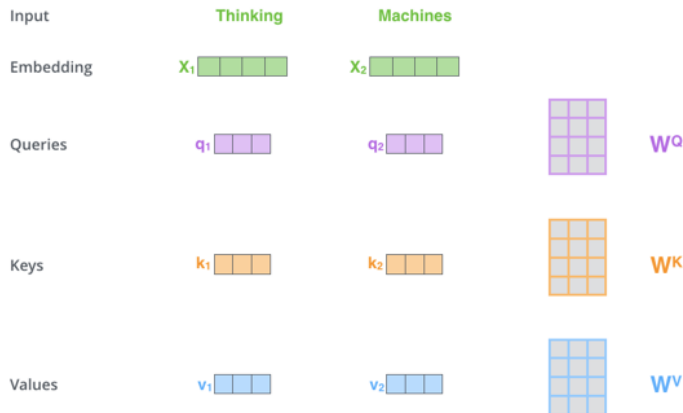


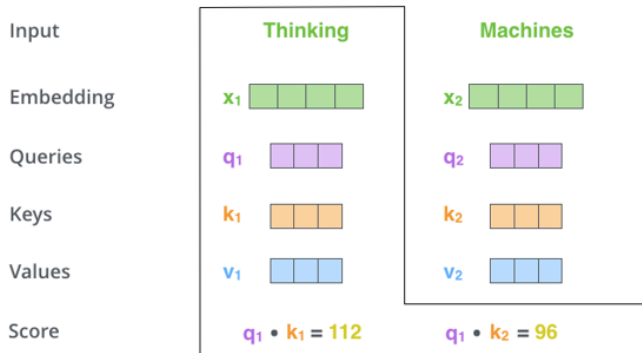
Figure: Multi-head Attention

Self-Attention



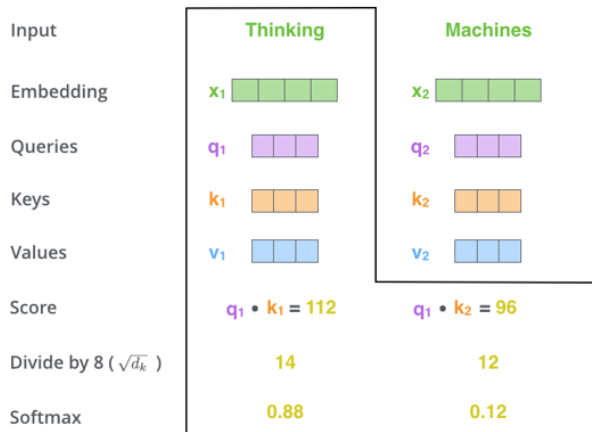
- Create three vectors for query, key, and value
- Multiplying x_1 by the W^Q weight matrix produces q_1 , and so on.

Self-Attention



- Mapping a query to keys (i.e., calculating the self-attention scores)

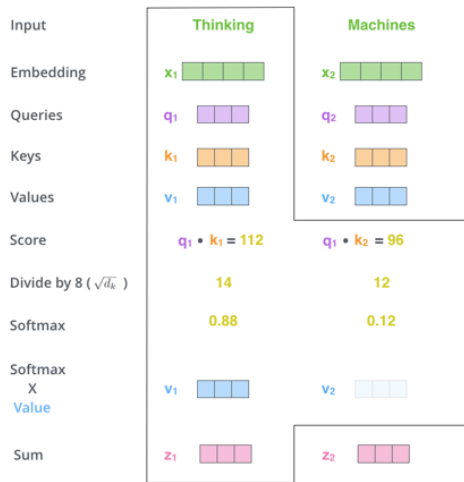
Self-Attention



- Normalize by the square root of the dimension of the key vectors
- Pass through a softmax operation

Self-Attention

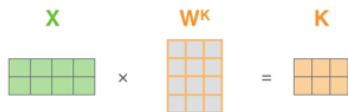
- Return self-attention score weighted value
- The accumulated weighted value vector (z 's) is the output of the self-attention layer at this position.



Self-Attention

- Matrix calculation by packing our embeddings into a matrix X , and multiplying it by the weight matrices we've trained (W 's)

$$X \times W^Q = Q$$


$$X \times W^K = K$$


$$X \times W^V = V$$


Self-Attention

$$\text{softmax}\left(\frac{\begin{matrix} \text{Q} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} \text{K}^T \\ \begin{array}{|c|} \hline \square \\ \hline \square \\ \hline \square \\ \hline \end{array} \end{matrix}}{\sqrt{d_k}}\right) \begin{matrix} \text{V} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix}$$
$$= \begin{matrix} \text{Z} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix}$$

- All the steps in one formula to calculate the outputs of the self-attention layer

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Self-Attention

1) This is our input sentence*

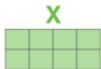
2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices


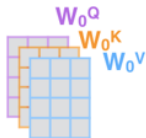
4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



Thinking
Machines



* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



...



...

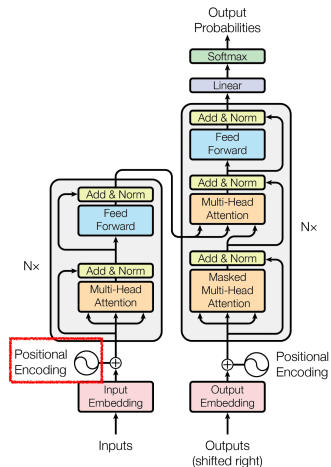


...



Positional Encoding

- The Transformer architecture is **missing** a way to account for **the order of the words** in the input sequence.
- Positional embeddings add positional information in the input sequence to each input embedding.
 - ▶ $PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{2i/d}}\right)$
 - ▶ $PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{2i/d}}\right)$



positional encoding

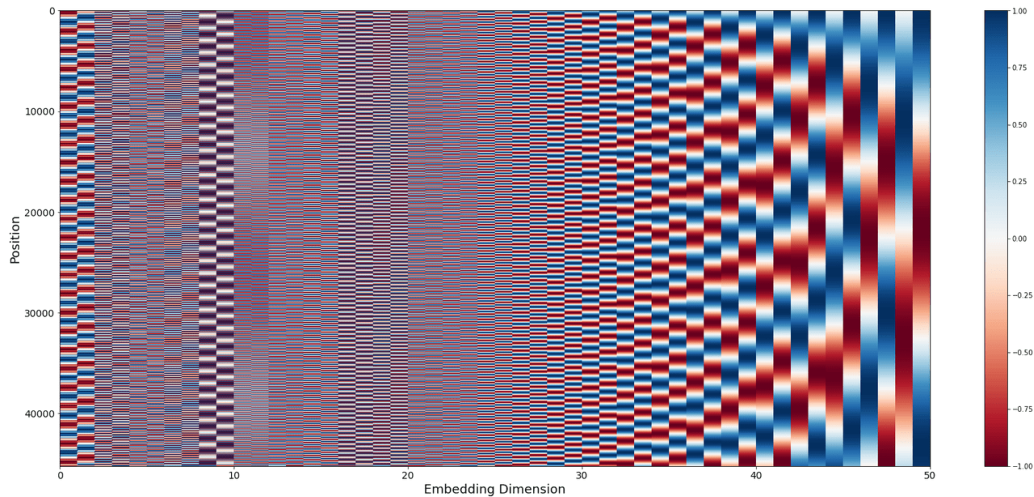
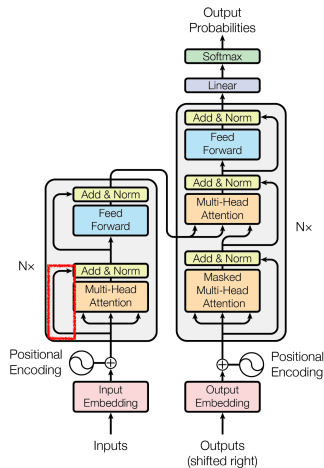


Figure: image by kemal erdem

Residual Connection

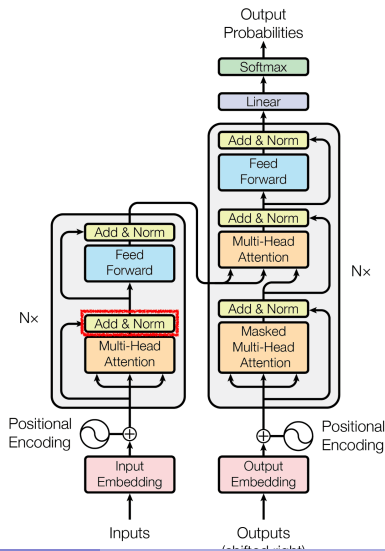
- Residual Connection
 - ▶ adds direct connection between the input layer and output layer,
 - ▶ allows gradients to flow through a network directly, without passing through non-linear activation functions.
 - ▶ helps to address the vanishing gradients problem during backpropagation.



Add and Normalization

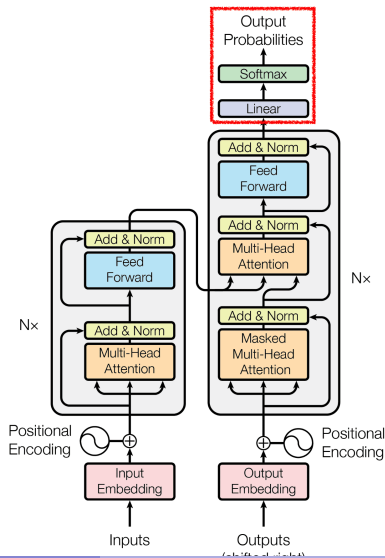
- Layer Normalization

- ▶ performed independently across features on each training example and for each layer.
- ▶ enables smoother gradients, faster training, and better generalization accuracy.

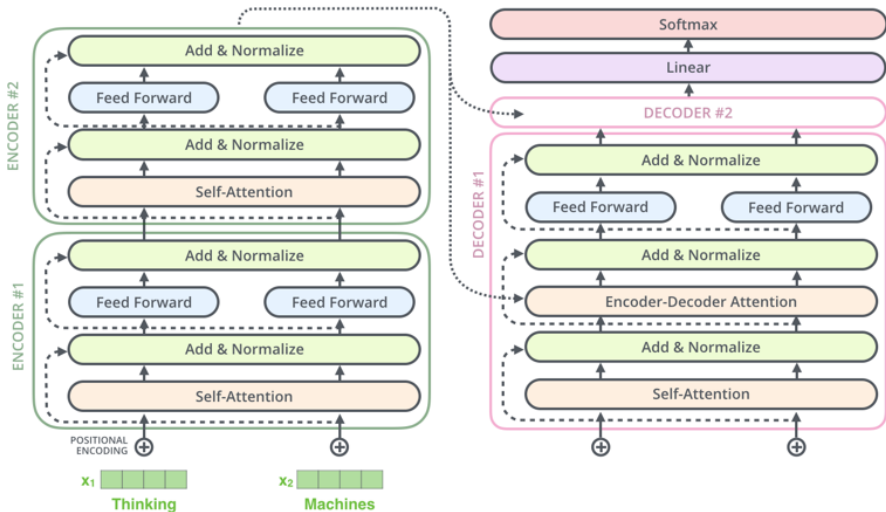


Prediction Layers

- The **final linear layer** outputs a much larger vector called a **logits vector**,
- and the following **softmax layer** converts it into log probabilities, which can be interpreted as a **probability distribution** in the target space associated with the downstream task.
 - ▶ If the task is language generation, the vector will be a probability distribution across the vocabulary.



Encoder-Decoder Attention



Topics

1 Attention Mechanism

2 Transformer Network

3 **Pre-trained Language Models**

- BERT
- GPT to ChatGPT

Language Modeling

- A **language model** (LM) is a probability distribution over sequences of token (e.g., words) variables, such that

$$P(s) = P(w_1, w_2, w_3, \dots w_n),$$

given a sequence $s = (w_1, w_2, w_3, \dots w_n)$.

- Typically, for optimization, we use the maximum likelihood estimation (MLE) with respect to the model parameters.

Language Modeling

- Optimization

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \prod_{i=1}^n P_{\theta}(w_i) \\ &= \arg \max_{\theta} \sum_{i=1}^n \log P_{\theta}(w_i) \\ &= \arg \max_{\theta} \sum_{i=1}^n \sum_{j=1}^{m_i} \log P_{\theta}(w_{i,j} | w_{i,1}, w_{i,2}, \dots, w_{i,j-1})\end{aligned}$$

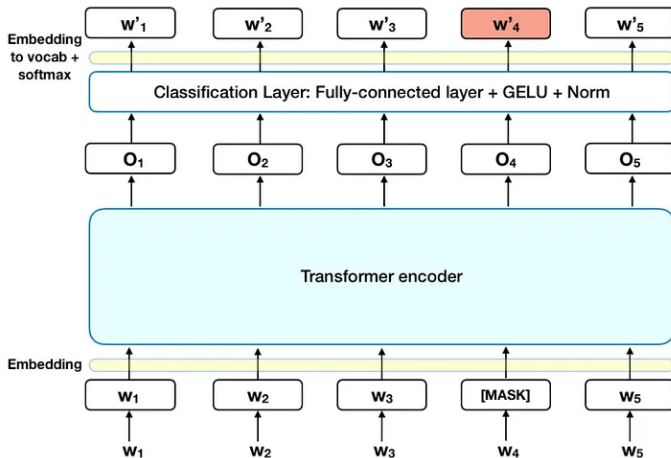
- where n is the number of examples and m_i is the sequence length.

BERT — Pre-trained Language Model


BERT (Bidirectional Encoder Representations from Transformers)

- A transformer (encoder) model trained using a large amount of unannotated and unstructured text data from the Internet.
- BERT learns the relationships between words in a language.
- BERT is trained on two specific tasks:
 - ▶ **Masked Language Modeling:** Predicting a missing word (masked tokens) in a sentence.
 - ▶ **Next Sentence Prediction:** Given two sentences A and B, predict if B comes after A.


BERT — Masked Language Model





BERT — Fine Tuning



(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



Topics

1 Attention Mechanism

2 Transformer Network

3 **Pre-trained Language Models**

- BERT
- GPT to ChatGPT

GPT-1

- GPT-1 (Generative Pre-Training) developed by OpenAI in 2018.
- They argued that there's no need of fine-tuning a pre-trained language model for specific natural language processing tasks (e.g., question-answering, document classification, etc.), and
- **autoregressive decoder model** is sufficient enough for most of the tasks.
- [paper](#) "Improving Language Understanding by Generative Pre-Training"

GPT-1 (Cont.)

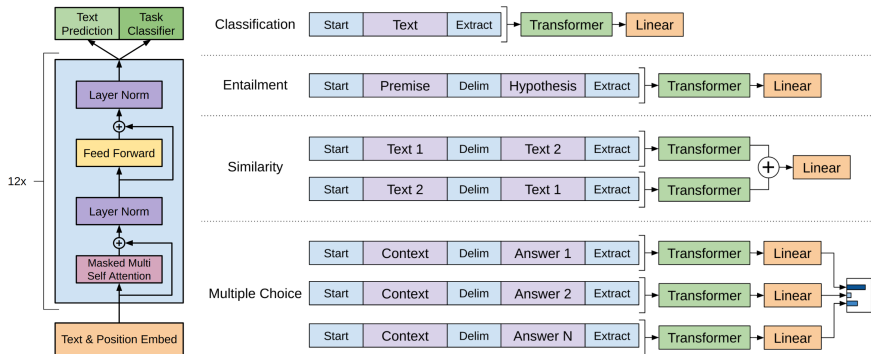


Figure 1: **(left)** Transformer architecture and training objectives used in this work. **(right)** Input transformations for fine-tuning on different tasks. We convert all structured inputs into token sequences to be processed by our pre-trained model, followed by a linear+softmax layer.

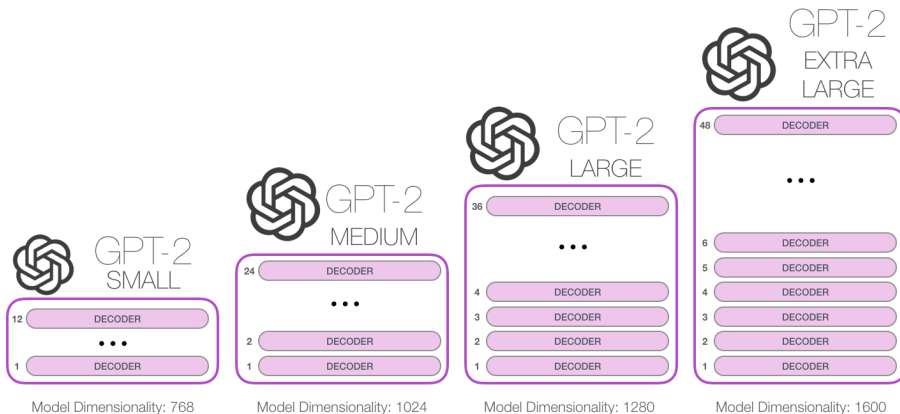
GPT-1 (Cont.)

- **GPT-1** is an autoregressive Transformer **decoder** model,
- whereas **BERT** is a non-autoregressive Transformer **encoder** model.
 - ▶ An autoregressive model can only pay attention to the previous tokens.
 - ▶ A non-autoregressive model can pay attention to tokens before and after itself bidirectionally.

GPT-2

- The second-generation of the GPT models released in 2019.
- Has more than one billion parameters (attributed to the collaboration of OpenAI and NVIDIA)
- Trained using an extremely large dataset, which might contain “annotated” examples for specific language tasks (e.g., QA or translation).
 - ▶ The extreme number of parameters and training examples enables the model *generating* answers, instead of *finding* answers.
- [paper](#) “Language Models are Unsupervised Multitask Learners”

GPT-2 (Cont.)



GPT-2 — Prompt-tuning

- **Prompt-tuning** is a technique to fine-tune large pre-trained language models to generate specific responses for a given prompt or context.
- This enables the model to generate more relevant and accurate responses for a given prompt.
- Useful for NLP certain tasks like chatbot, question-answering system, language translation.
- GPT-2 started to use **prompt-tuning** to guide the model in training for question answering, summarization, translation, and so on.
 - ▶ We can give the model a prompt to have it speak about a certain topic.
 - ▶ E.g., **recite the first law** \$ **A robot may not injure a human being ...**

GPT-3

- The third-generation of the GPT models released in 2020.
- Even larger in size (175 billion model parameters)
- GPT-3 is trained to read and understand the context and repeat generating answer in the same manner.
- [paper](#) “Language Models are Unsupervised Multitask Learners”

GPT-3 Example

Prompt

I am a highly intelligent question answering bot. If you ask me a question that is rooted in truth, I will give you the answer. If you ask me a question that is nonsense, trickery, or has no clear answer, I will respond with "Unknown".

Q: What is human life expectancy in the United States?

A: Human life expectancy in the United States is 78 years.

Q: Who was president of the United States in 1955?

A: Dwight D. Eisenhower was president of the United States in 1955.

Q: Which party did he belong to?

A: He belonged to the Republican Party.

Q: What is the square root of banana?

A: Unknown

Q: How does a telescope work?

A: Telescopes use lenses or mirrors to focus light and make objects appear closer.

Q: Where were the 1992 Olympics held?

A: The 1992 Olympics were held in Barcelona, Spain.

Q: How many squigs are in a bonk?

A: Unknown

Settings

Engine	text-davinci-003
Max tokens	100
Temperature	0
Top p	1
Frequency penalty	0.0
Presence penalty	0.0
Stop sequence	\n

InstructGPT

- Issues with the previous GPT models
 - ▶ The models are not trained to follow human instructions.
 - ★ It needs to be provided with contextualized demonstration.
 - ★ Triggers are engineered, e.g. "A:" or "Answer:"
 - ★ Questions without the context is difficult to be answered.
- InstructGPT developed to address these issues in 2022. (paper, "[Training Language Models to Follow Instructions with Human Feedback](#)")

InstructGPT (Cont.)

Training the InstructGPT model

- Collect demonstration data of desired output is created by human beings. Use this data to fine-tune GPT-3 using supervised learning.
- Collect comparison data: GPT-3 generates outputs to a prompt, and a human labeler ranks the multiple outputs which is then used to train the model using reinforcement learning.

InstructGPT — Training


Step 1

Collect demonstration data, and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.



This data is used to fine-tune GPT-3 with supervised learning.



Step 2

Collect comparison data, and train a reward model.


A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.



Step 3

Optimize a policy against the reward model using reinforcement learning.

A new prompt is sampled from the dataset.




The policy generates an output.



The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.



ChatGPT

- ChatGPT and InstructGPT are essentially the same model, except that for ChatGPT,
 - ▶ Different demonstration examples are used to train for better conversation between the users and the bot.
 - ▶ E.g., Examples include dialogues created by human.
- This interactive information system can replace conventional search engines in many different ways.